

Aditya Kulkarni

[linkedin.com/in/aditya-kulkarni-355b81217](https://www.linkedin.com/in/aditya-kulkarni-355b81217)

github.com/Aditya-k24 kulkarniaditya.com

 (281) 876-7066  adityakulkarnius@gmail.com 

June 21, 2026

Hiring Team, Intel Corporation
Software Engineer – Agent Harness (JR0284869)
Santa Clara, CA

Dear Hiring Team,

I'm applying for the Software Engineer – Agent Harness role. I've spent the last year building the exact stack Intel is describing: MCP servers for tool schemas and context injection, Temporal.io for durable agent orchestration, and a Kubernetes-native inference platform that routes across multiple model backends.

What I bring to this role:

- **Agent harness engineering, end-to-end.** At Isomer AI, I built MCP servers in both Python and TypeScript that power configurable multi-model LLM agent pipelines (Claude, GPT-4o, Gemini) with typed field extraction schemas, context injection, and Playwright browser automation – deployed to 30+ enterprise customers. I also engineered the Temporal.io workflow engine driving these pipelines: 3-state versioning (Draft → Live → Archived), canary traffic splits, QA sampling, and CRITICAL/HIGH/STANDARD risk signal detection.
- **LLM inference infrastructure.** CortexQ is a Kubernetes-native LLM inference platform I built independently: a FastAPI router with a 3-state circuit breaker and EWMA latency-aware load balancing across Claude, GPT-4o, and Gemini backends; a KEDA autoscaler that scales model pods from 0 to N on Redis queue depth; and a Python Kopf CRD operator (LLMDeployment) that reconciles Deployments, Services, and ScaledObjects. This maps directly to Intel's local+cloud hybrid intelligence architecture – the same scale-to-zero and dynamic routing logic applies to edge inference.
- **Model benchmarking and evaluation.** My LLM Eval Harness compares Claude Sonnet 4.6, GPT-4o, and Gemini 2.5 Pro on a 2,000-example frozen test set across grounded QA, multi-hop QA, and claim verification – tracking ExactMatch, TokenF1, grounded hallucination rate, p50/p95 latency, and cost/call with paired bootstrap CIs and McNemar significance testing. This is the kind of rigorous, vendor-neutral evaluation Intel's agent harness needs.

I'm graduating from NC State's MS CS program in May 2026 and am ready to relocate to Santa Clara immediately. I would welcome the chance to discuss how my background maps to Intel's agent harness work.

Sincerely,

Aditya Kulkarni