

June 22, 2026

Hiring Team, Microsoft
Software Engineer, AI Development Acceleration Program (MAIDAP)
Redmond, WA

Dear Hiring Team,

I'm applying for the Software Engineer role in the AI Development Acceleration Program. MAIDAP's mission, providing AI as a service to engineering teams across Microsoft, is exactly the work I've spent the last year doing: building LLM infrastructure that other engineers plug into and ship on.

What I bring to MAIDAP:

- **AI as a service, in production.** At Isomer AI, I built MCP servers in Python and TypeScript that expose typed tools and context to multi-model LLM agents (Claude, GPT-4o, Gemini), and engineered the Temporal.io workflow engine that runs them durably: versioned releases, canary splits, QA sampling, and automated risk detection. Product teams consumed this as a platform, not a one-off.
- **Inference infrastructure, end-to-end.** CortexQ is a Kubernetes-native LLM inference platform I built independently: a FastAPI router with a circuit breaker and EWMA latency-aware load balancing across model backends, a KEDA autoscaler that scales pods to zero when idle, and a Python Kopf operator that reconciles the whole deployment. This is the local-plus-cloud, multi-team serving pattern MAIDAP works on.
- **Rigorous, vendor-neutral evaluation.** My LLM Eval Harness benchmarks Claude Sonnet 4.6, GPT-4o, and Gemini 2.5 Pro on a frozen 2,000-example set, tracking ExactMatch, TokenF1, hallucination rate, latency, and cost with bootstrap confidence intervals and McNemar tests, the kind of measurement discipline a shared AI platform depends on.

I'm completing my MS in Computer Science at NC State in May 2026 and am ready to start full-time in July 2026 and relocate to Redmond. I'd welcome the chance to bring this work to MAIDAP's teams.

Sincerely,

Aditya Kulkarni