

Aditya Kulkarni

☎ (281) 876-7066 | ✉ adityakulkarnius@gmail.com

🌐 [linkedin.com/in/aditya-kulkarni-355b81217](https://www.linkedin.com/in/aditya-kulkarni-355b81217) | 🏠 github.com/Aditya-k24

June 12, 2026

To the 7-Eleven AI Engineering Team,

The AI SDLC team at 7-Eleven is solving a problem I find genuinely interesting: applying language models not to end users, but to the engineers building software at scale. Code generation, test creation, RAG-based knowledge search, and CI/CD automation across an organization with 86,000+ global locations is a different kind of reliability challenge than a typical AI product. I am completing my Master's in Computer Science at NC State and want to build these tools.

- › **LLM Evaluation — Built From Scratch:** I built a vendor-neutral evaluation harness in Python benchmarking Claude, GPT-4o, and Gemini across 2,000 frozen test examples: exact match, token F1, hallucination rate, p50/p95 latency, cost-per-call, and McNemar significance tests. Evaluating AI output quality and reliability is the core responsibility of this team, and I have built the tooling to do it rigorously.
- › **AI Into CI/CD Pipelines:** CortexQ ships multi-provider LLM routing as a Kubernetes platform with GitHub Actions CI, Helm + ArgoCD GitOps, and OpenTelemetry observability end-to-end. Integrating AI capabilities into engineering platforms is exactly the scope described in this role.
- › **LangChain + RAG + Workflow Orchestration:** At Isomer AI, I shipped production LangChain pipelines across Claude, GPT-4o, and Gemini with Temporal.io workflow orchestration — the same stack the JD calls out: LangChain, RAG, tool orchestration, AI workflow automation.
- › **Cross-Functional Delivery:** At Isomer AI, I collaborated daily with product managers, designers, and engineers across 923 commits and 7 production deployments. At Mphasis, I worked with health-care data teams and led model validation and red-teaming sessions. 7-Eleven's SDLC team requires the same cross-functional communication, and I have practiced it in production.

I want to build AI tools that make engineers faster and software more reliable. 7-Eleven's scale makes that problem real and meaningful. I would welcome the chance to discuss.

Sincerely,

Aditya Kulkarni