

Aditya Kulkarni

📍 Sunnyvale, CA | 📞 (281) 876-7066 | ✉️ adityakulkarnius@gmail.com
🌐 linkedin.com/in/aditya-kulkarni-355b81217 | 🐙 github.com/Aditya-k24

June 14, 2026

To the NGSiem Agentic AI Team at CrowdStrike,

CrowdStrike processes almost 3 trillion security events per day. Building intelligent agentic systems on top of that data—retrieval pipelines, ML models, and AI agents that actually work in adversarial, high-stakes conditions—is the kind of engineering I want to build my career on. The NGSiem Agentic AI Engineer I role sits at exactly that intersection.

- **Agentic AI in Production:** At **Isomer AI**, I shipped 10+ production AI features into enterprise workflows—multi-model LLM pipelines (Claude, GPT-4o, Gemini), Temporal.io agentic workflows with MCP tooling, and full Auth0 RBAC—across 923 commits and 7 deployments. I owned the backend through incidents and SLA commitments, not just initial shipping. My **AlgoPulse** project scaled an event-driven AI orchestration platform to 2,000 async inference requests/second using Temporal.io and Kafka for fault-tolerant state synchronization.
- **Search, Retrieval, and ML Pipelines:** I built **VectorLift**, a production-grade semantic search system on MS MARCO combining BM25 (Elasticsearch), fine-tuned bi-encoder retrieval, and cross-encoder re-ranking—evaluated across 6 pipeline configurations with NDCG@10, MRR@10, and MAP, with a Kafka-driven indexing pipeline feeding both Elasticsearch and Qdrant. This is the retrieval architecture that NGSiem needs: fast lexical search, dense retrieval, and reranking at query time.
- **Distributed Systems and Python:** My **CortexQ** project is a Kubernetes-native LLM inference platform in Python: FastAPI router with EWMA latency-aware load balancing, a Kopf CRD operator reconciling KEDA ScaledObjects, circuit breaker logic, and an OpenTelemetry + Prometheus + Grafana observability stack—load tested to 500 VUs and validated with 88 pytest integration tests. At **Capgemini**, I built Kafka + Spark + Hadoop pipelines that cut data-processing latency by 20%.
- **Cybersecurity Background:** My Springer-published research on *Computer Vision-Based Cybersecurity Threat Detection* (icSoftComp 2023) is directly relevant to CrowdStrike's mission. I understand threat modeling and adversarial data, not just AI engineering—a combination that matters on the NGSiem team where the models run on real attack telemetry.

CrowdStrike's problem space—AI agents reasoning over petabyte-scale security telemetry in real time—is one of the hardest engineering challenges in industry. I want to work on it from day one. I would welcome the opportunity to discuss how my background fits the NGSiem team.

Sincerely,

Aditya Kulkarni