

Senior SDE — AI Data Platform (FDE), Oracle

June 20, 2026

Oracle's FDE AI Data Platform team is building the infrastructure that makes generative AI and real-time analytics work at OCI scale. I have built both sides of that stack — LLM serving infrastructure and AI data platform engineering — in production.

Why I am a strong fit:

- **LLM serving platform at infrastructure depth.** KubeServe: Kubernetes-native LLM inference platform with FastAPI router, KEDA autoscaler (scale-to-zero to N replicas), Kopf CRD operator managing multi-provider LLM backends, circuit breaker + EWMA load balancing, and OTel + Prometheus + Grafana full observability. This is the "power generative AI and LLMs at scale" work the FDE team does — I built it from scratch.
- **Real-time analytics data platform.** VectorLift: Kafka-driven indexing pipeline into Elasticsearch and Qdrant vector store, FastAPI search API, Streamlit analytics dashboard benchmarking 6 retrieval pipeline configurations with NDCG@10, MRR@10, and MAP. Built to understand which data platform decisions move the metrics that matter.
- **Enterprise customer deployments (FDE context).** At Isomer AI I was the engineer who shipped and owned 30+ per-tenant production deployments for enterprise insurance carriers — each with its own Fly.io + Supabase + Temporal Cloud stack. The FDE role requires shipping reliable AI systems directly into customer environments; that is exactly what I did at Isomer.
- **AI evaluation and observability across providers.** LLM Eval Harness: benchmarked Claude, GPT-4o, and Gemini across 2,000 frozen labeled examples tracking hallucination rate, p50/p95 latency, cost-per-call, and 6 metrics with McNemar significance tests. Data-driven LLM selection is the foundation of any production AI data platform.

OCI is building the AI infrastructure layer that the next decade runs on. I want to help build the data platform that makes AI systems at that scale actually work.

Aditya Kulkarniadityakulkarnius@gmail.com | (281)-876-7066