

Aditya Kulkarni

☎ (281) 876-7066 | ✉ adityakulkarnius@gmail.com

🌐 [linkedin.com/in/adityakulkarni244](https://www.linkedin.com/in/adityakulkarni244) | 🏠 github.com/Aditya-k24

June 20, 2026

To the Casca Team,

Legacy banking software exists to protect the bank, not to serve the people using it. Every approval workflow, every document review, every risk assessment that a human does manually today is a process waiting to be automated with AI. Casca is doing exactly that — replacing the manual layer with AI-native infrastructure from the ground up. The Forward Deployed Engineer role is about sitting at the seam between Casca's technology and banking customers who need it to actually work in production. That is where I want to be.

- **AI for Regulated, High-Stakes Industries, Shipped:** At **Isomer AI**, I shipped a per-tenant AI workflow platform for insurance carriers — Temporal.io durable orchestration, LLM risk signal detection across Claude/GPT-4o/Gemini, Auth0 RBAC, MCP servers, 30+ deployments. Insurance is structurally similar to banking: regulated, high-stakes, full of manual review workflows that AI can automate. I built that system end to end for enterprise clients.
- **LangChain + Vercel AI SDK, Production:** The stack Casca lists — React, TypeScript, Next.js, Python, LangChain — is exactly what I build with. LangChain is in my skills and in production at Isomer. I have used it to build RAG pipelines, tool-calling agents, and multi-model routing. Not demos; production features with real users.
- **Evaluation Rigor for AI You Can Trust:** Banking AI cannot hallucinate. I built an LLM evaluation harness from scratch — 2,000 frozen test examples, hallucination rate, p50/p95 latency, cost-per-call, McNemar significance tests across Claude, GPT-4o, and Gemini. The rigor behind knowing when your AI is wrong is what separates deployable banking software from a prototype.
- **Agentic Systems at Production Scale:** **AlgoPulse** is an agentic AI orchestration platform I built on Temporal.io and Kafka — 2,000 async inference requests per second, sub-200ms delivery, event-driven SSE replacing polling with 60% fewer active server connections. The core skill of an FDE is building AI systems fast and making them hold. I have done that independently.

Casca is compressing decades of banking infrastructure change into a few years. I want to help close the gap between what the technology can do and what banking customers are actually running in production. I would be glad to discuss.

Sincerely,

Aditya Kulkarni