

Aditya Kulkarni

☎ (281) 876-7066 | ✉ adityakulkarnius@gmail.com

🌐 linkedin.com/in/adityakulkarni244 | 🏠 github.com/Aditya-k24

June 21, 2026

To the Glean Engineering Team,

Agents fail at runtime, not at design time. The reliability, latency, and observability gaps in production agentic systems are exactly the problems I have been solving — circuit breakers that prevent cascading LLM failures, OTel traces that make multi-turn tool invocations debuggable, and Temporal.io orchestration that gives agents durable, recoverable execution. Glean’s Agentic Runtime team is building this infrastructure at enterprise scale. I want to help build it.

- **Distributed Agentic Runtime, Built from First Principles:** KubeServe is a Kubernetes-native LLM inference platform I built: FastAPI router, KEDA autoscaler, Kopf CRD operator for multi-provider LLM serving, circuit breaker with EWMA load balancing, and full OTel + Prometheus + Grafana observability. This is the same fault-isolation, retry, and backpressure architecture that Glean’s runtime needs at scale.
- **Multi-Turn Orchestration with Durable Execution:** At Isomer AI, I built production agentic workflows on Temporal.io: per-tenant durable orchestration, LLM risk signal pipelines across Claude/GPT-4o/Gemini, and MCP server tool integrations. In AlgoPulse, I extended this to Kafka-backed multi-turn reasoning at 2,000 async inference requests per second with sub-100ms streaming delivery. Temporal.io durable execution, tool invocation, retry, and termination logic — I have implemented all of it in production.
- **Multi-LLM Provider Integration and Eval Frameworks:** LLM Eval Harness is a vendor-neutral Python evaluation system I built comparing Claude, GPT-4o, and Gemini across 2,000 frozen labeled examples — hallucination rate, p50/p95 latency, cost-per-call, McNemar significance tests. The role calls out “integrate with leading LLM providers and internal evaluation frameworks” — this is exactly what I built.
- **Production Observability and SLOs:** Across KubeServe and Isomer, I instrumented agent pipelines end-to-end with OpenTelemetry traces, Prometheus metrics, and Grafana dashboards. I have set up latency SLOs, tail behavior monitoring, and on-call playbooks for agentic systems.

Glean’s agentic runtime is the infrastructure that makes enterprise AI reliable, not just impressive. I want to help build it to the standard enterprise customers demand. I would be glad to discuss.

Sincerely,

Aditya Kulkarni