

Aditya Kulkarni

☎ (281) 876-7066 | ✉ adityakulkarnius@gmail.com

🌐 [linkedin.com/in/aditya-kulkarni-355b81217](https://www.linkedin.com/in/aditya-kulkarni-355b81217) | 🐙 github.com/Aditya-k24

June 19, 2026

To the Intel AI Engineering Team,

Agentic AI frameworks are multiplying faster than the tooling to evaluate them. The gap between deploying an agent framework and knowing whether it is actually reliable, regressing, or cost-efficient is where most teams are flying blind. Building the evaluation infrastructure to close that gap — deployment, observability, regression pipelines, and dashboards — is exactly the work I have been doing independently, and it is exactly what this role builds at Intel's scale.

- **LLM Evaluation Infrastructure, Built from Scratch:** I built a vendor-neutral LLM evaluation harness in Python — 2,000 frozen labeled examples, hallucination rate, p50/p95 latency, cost-per-call, McNemar significance tests across Claude, GPT-4o, and Gemini. This directly maps to the role's mandate of “automated evaluation pipelines for performance, reliability, and regression detection.”
- **Scalable AI Deployment and Orchestration:** **CortexQ** is a Kubernetes-native LLM inference platform I built: FastAPI router, KEDA autoscaler, Kopf CRD operator for multi-provider serving, circuit breaker, EWMA load balancing, Helm + ArgoCD GitOps, and full OpenTelemetry + Prometheus + Grafana observability. I know what it takes to build deployment infrastructure that holds under real load.
- **Agentic Framework Integration at Scale:** **AlgoPulse** is a production agentic AI orchestration platform on Temporal.io and Kafka — 2,000 async inference requests per second, sub-200ms SSE delivery, multi-step reasoning pipelines. I have designed, debugged, and shipped the kind of agentic systems this role exists to evaluate.
- **Multi-Model API Integration and Prompt Engineering:** At **Isomer AI**, I shipped production LLM pipelines integrating Claude (Anthropic), GPT-4o, and Gemini with tool-calling, RAG, and Temporal.io agentic workflows. The role specifically calls out Anthropic API integrations — I have used them in production, not just in demos.

Intel is building the infrastructure layer that the next generation of AI runs on. The evaluation ecosystem for agentic AI is a foundational part of that. I want to help build it. I would be glad to discuss.

Sincerely,

Aditya Kulkarni