

Aditya Kulkarni

📍 Raleigh, NC | 📞 (281) 876-7066 | ✉️ adityakulkarnius@gmail.com
🌐 linkedin.com/in/aditya-kulkarni-355b81217 | 🐙 github.com/Aditya-k24

June 10, 2026

To the AWS Data Services Engineering Team,

AWS operates the database, storage, and analytics infrastructure that millions of applications depend on every day. The scale at which AWS Data Services works—billions of requests, petabyte-scale storage, sub-millisecond latency requirements—is a class of engineering problem I want to work on. I bring production experience in distributed systems, data pipelines, cloud-native architecture on AWS, and a track record of shipping software that runs reliably under load.

- **Distributed Systems and Production Ownership:** At **Isomer AI**, I own the full production backend—PostgreSQL, Redis, Docker, Terraform—managing CI/CD, resolving on-call incidents, and maintaining SLA commitments across enterprise environments. I built **CortexQ**, a Kubernetes-native distributed service platform with KEDA autoscaling, circuit breaker logic, EWMA load balancing, and OpenTelemetry observability, load tested to 500 concurrent users with 88 integration tests validating correctness.
- **Data Pipelines and Analytics Systems:** At **Capgemini**, I built distributed Kafka streaming pipelines for real-time ingestion and Spark + Hadoop batch processing that cut ticket-processing latency by 20%. I built an **LLM Eval Harness** using DuckDB and Parquet to store and query evaluation results across 2,000 labeled examples—tracking latency, cost, and quality metrics with Bootstrap confidence intervals and McNemar significance tests.
- **AWS and Cloud-Native Architecture:** My production experience spans AWS and GCP: automated CI/CD deployments at **Mphasis**, Grafana monitoring dashboards, containerized microservice deployments, and infrastructure-as-code with Terraform. AWS is the platform I work on and the platform AWS Data Services is built on.
- **Internship Track Record Across Three Companies:** **Isomer AI** (1 year production), **Mphasis** (HP Company), and **Capgemini**—three distinct engineering environments spanning AI products, enterprise software, and large-scale data systems. Each built my instinct for ownership, operational reliability, and shipping code that teams depend on.

AWS Data Services is the infrastructure layer that the rest of the cloud runs on. That is exactly the kind of foundational engineering challenge I want to work on. I would be glad to discuss how my experience maps to your team's needs.

Sincerely,

Aditya Kulkarni