

Employment

Product Engineer Intern	Isomer AI	May 2025 – May 2026
- Shipped 10+ production AI features: multi-model LLM pipelines (Claude, GPT-4o, Gemini) with tool calling and structured outputs, Temporal.io agentic workflows, and Auth0 RBAC; 923 commits, 7 deployments. - Own backend stack (PostgreSQL, Redis, Docker, Terraform, Sentry APM); maintain CI/CD and resolve production incidents.		
Graduate Research Assistant	NC State University	April 2025 – Present
- Deploy cloud genomic platform on GCP (Cloud Run) with JBrowse 2 genome browser; support 100+ researchers, 500GB+ datasets. - Improve ML model accuracy by 12% via hyperparameter tuning and time-series analysis on biological sensor data.		
Software Developer Intern	Mphasis (HP Company)	Jun 2024 – Aug 2024
- Built 5+ AI workflows including health claim automation (OCR, NLP, ML); improved document processing accuracy by 25%. - Automated AWS and GCP deployments via CI/CD; integrated Grafana monitoring and led model red-teaming for reliability.		
AI/ML Intern	Capgemini	Jun 2023 – Aug 2023
- Led SDLC for data-driven ticket analysis system using NLP classification; reduced average resolution time by 15%. - Built Kafka + Spark + Hadoop pipelines for batch and real-time data ingestion; cut ticket-processing latency by 20%.		

Education

Raleigh, NC	North Carolina State University	Aug 2024 – May 2026
- Master of Science in Computer Science - Courses: Artificial Intelligence, Operating Systems, Algorithms, Data Science, Object Oriented Development, Software Engineering.		
Mumbai, India	Mumbai University	Aug 2020 – May 2024
Bachelor of Technology in Computer Science; CGPA: 9.32/10. Courses: Data Warehousing, Big Data Infra, DS & Algorithms, Machine Learning, Blockchain.		

Projects

LLM Eval Harness	AI Data Analysis Platform	2025
- Vendor-neutral evaluation harness comparing Claude, GPT-4o, Gemini across 2,000 frozen examples; tracked ExactMatch, TokenF1, hallucination rate, abstention accuracy, p50/p95 latency, and cost-per-call. - Bootstrap CIs, McNemar test, Holm correction across 3 repeated runs per model; DuckDB + Parquet + Streamlit dashboard.		
Glance AI	Fintech/Healthcare Privacy RAG	2025
- Per-tenant RAG for fintech/healthcare with vector isolation preventing cross-tenant data leakage and hallucinations. - Deterministic retrieval architecture for regulated environments; factual LLM outputs, no context bleed between users.		
AlgoPulse	Full-Stack Agentic AI Platform	2025
- Built agentic AI orchestration platform on Temporal.io and Kafka supporting multi-step reasoning pipelines; sustained 2,000 async inference requests per second under sustained production load. - Replaced API polling with Supabase Edge Functions for real-time AI streaming; eliminated state loss under peak traffic.		
CortexQ	Kubernetes LLM Inference Platform	2026
- Built Kubernetes-native LLM inference platform: FastAPI router, KEDA autoscaler (scale-to-zero to N replicas), Kopf CRD operator for multi-provider LLM serving; circuit breaker + EWMA load balancing. - Instrumented OTel + Prometheus + Grafana observability; Helm + ArgoCD GitOps; k6 to 500 VUs; 88 pytest tests.		

Languages and Technologies

- Languages:** Python, TypeScript, JavaScript, Java, Ruby, SQL, Bash
- AI/LLM & Data:** Claude/GPT-4o/Gemini/Vertex AI APIs, RAG, LLM evaluation, DuckDB, Parquet, Streamlit, pandas
- Backend & APIs:** FastAPI, Node.js, NestJS, PostgreSQL, Redis, REST APIs, Kafka, Temporal.io, Auth0
- Cloud & DevOps:** GCP, AWS, Docker, Kubernetes, Terraform, GitHub Actions, OpenTelemetry, Prometheus, Grafana

Additional Experience and Awards

- IEEE Published:** Co-authored paper on CV-Based Cybersecurity Threat Detection (Springer icSoftComp 2023); presented to 500+ attendees.
- National Hackathon Winner:** Won or ranked Top 10 in 6+ national hackathons spanning ML, systems, and developer tools; built 3 complete MVPs end-to-end within 48-hour sprints, each with a working demo.
- Lead at Google Developers Student Club:** Led 4+ tech events (hackathons, conferences) reaching 500+ students.
- Technical Writing:** Published on LLM model collapse, AI agent security, audio AI at kulkarniaditya.com/blog